



# Python threads: Dive into GIL!

PyCon India 2011  
Pune, Sept 16-18

Vishal Kanaujia & Chetan Giridhar

# Python: A multithreading example

```
__author__ = "Chetan Giridhar, Vishal Kanaujia"
__date__ = "$Aug 31, 2011 09:37:00$"

from datetime import datetime
import threading
import random

def cpu(n):
    while n>0:
        n-=1
        (random.uniform(1, 10000))/(random.uniform(1, 100000))
        (random.uniform(1, 10000))*(random.uniform(1, 100000))

iterations = 12000000

startTime = datetime.now()

thread1 = threading.Thread(target=cpu, args=(iterations,))
thread2 = threading.Thread(target=cpu, args=(iterations,))


thread1.start()
thread2.start()

thread1.join()
thread2.join()

endTime = datetime.now()


diffTime = endTime - startTime
print diffTime
```

# Setting up the context!!



Hmm...My threads  
should be twice as faster  
on dual cores!


**Execution Time**



Python v2.7	Execution Time
Single Core	74 s
Dual Core	116 s

**57 % dip in Execution Time on dual core!!**

# Getting into the problem space!!



Python  
v2.7

GIL

# Python Threads

- Real system threads (POSIX/ Windows threads)
- Python VM has no intelligence of thread management
  - No thread priorities, pre-emption
- Native operative system supervises thread scheduling
- Python interpreter just does the per-thread bookkeeping.

# What's wrong with Py threads?

- Each 'running' thread requires exclusive access to data structures in Python interpreter
- Global interpreter lock (GIL) provides the synchronization (bookkeeping)
- GIL is necessary, mainly because CPython's memory management is *not* thread-safe.

# GIL: Code Details

- A thread create request in Python is just a `pthread_create()` call
- The function `Py_Initialize()` creates the GIL
- GIL is simply a synchronization primitive, and can be implemented with a semaphore/ mutex.

*../Python/ceval.c*

*static PyThread\_type\_lock interpreter\_lock = 0; /\* This is the GIL \*/*

- A “runnable” thread acquires this lock and start execution

# GIL Management

- How does Python manages GIL?
  - Python interpreter *regularly* performs a check on the running thread
  - Accomplishes thread switching and signal handling
- What is a “check”?
  - A counter of ticks; ticks decrement as a thread runs
  - A tick maps to a Python VM’s byte-code instructions
  - Check interval can be set with `sys.setcheckinterval(interval)`
  - Check dictate CPU time-slice available to a thread





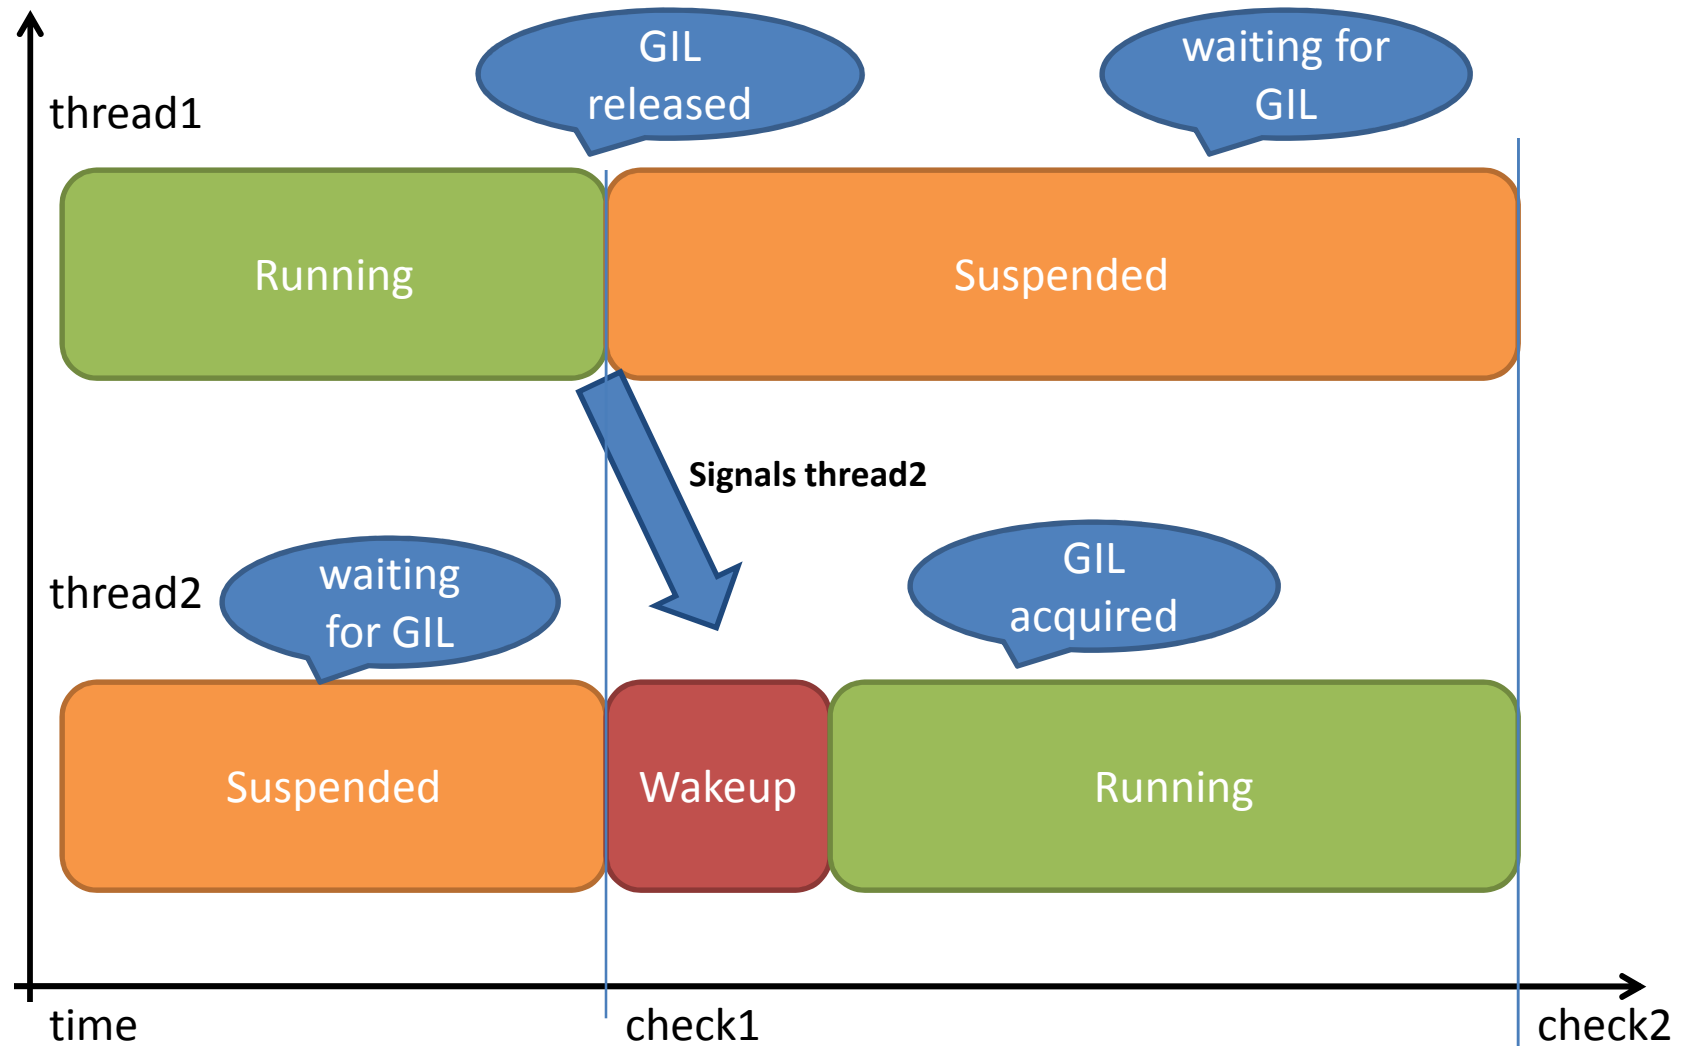
# GIL Management: Implementation

- Involves two global variables:
  - `PyAPI_DATA(volatile int) _Py_Ticker;`
  - `PyAPI_DATA(int) _Py_CheckInterval;`
- As soon as ticks reach zero:
  - Refill the ticker
  - active thread ***releases the GIL***
  - *Signals sleeping threads to wake up*
  - *Everyone competes for GIL*

File: ../ceval.c

```
PyEval_EvalFrameEx () {  
    --_Py_Ticker; /* Decrement ticker */  
    /* Refill it */  
    _Py_Ticker = _Py_CheckInterval;  
    if (interpreter_lock) {  
        PyThread_release_lock(interpreter_lock);  
  
        /* Other threads may run now */  
  
        PyThread_acquire_lock(interpreter_lock, 1);  
    }  
}
```

# Two CPU bound threads on single core machine




# GIL impact

- There is considerable time lag with
  - Communication (Signaling)
  - Thread wake-up
  - GIL acquisition
- GIL Management: Independent of host/native OS scheduling
- **Result**
  - Significant overhead
  - Thread waits if GIL is unavailable
  - Threads run sequentially, rather than concurrently

Try Ctrl + C. Does it stop execution?

# Curious case of multicore system




- Conflicting goals of OS scheduler and Python interpreter
- Host OS can schedule threads concurrently on multi-core
- GIL battle

# The ‘Priority inversion’


- In a [CPU, I/O]-bound mixed application, I/O bound thread may starve!
- “cache-hotness” may influence the new GIL owner; usually the recent owner!
- Preferring CPU thread over I/O thread
- Python presents a *priority inversion* on multi-core systems.

# Understanding the new story!!



Python  
v3.2

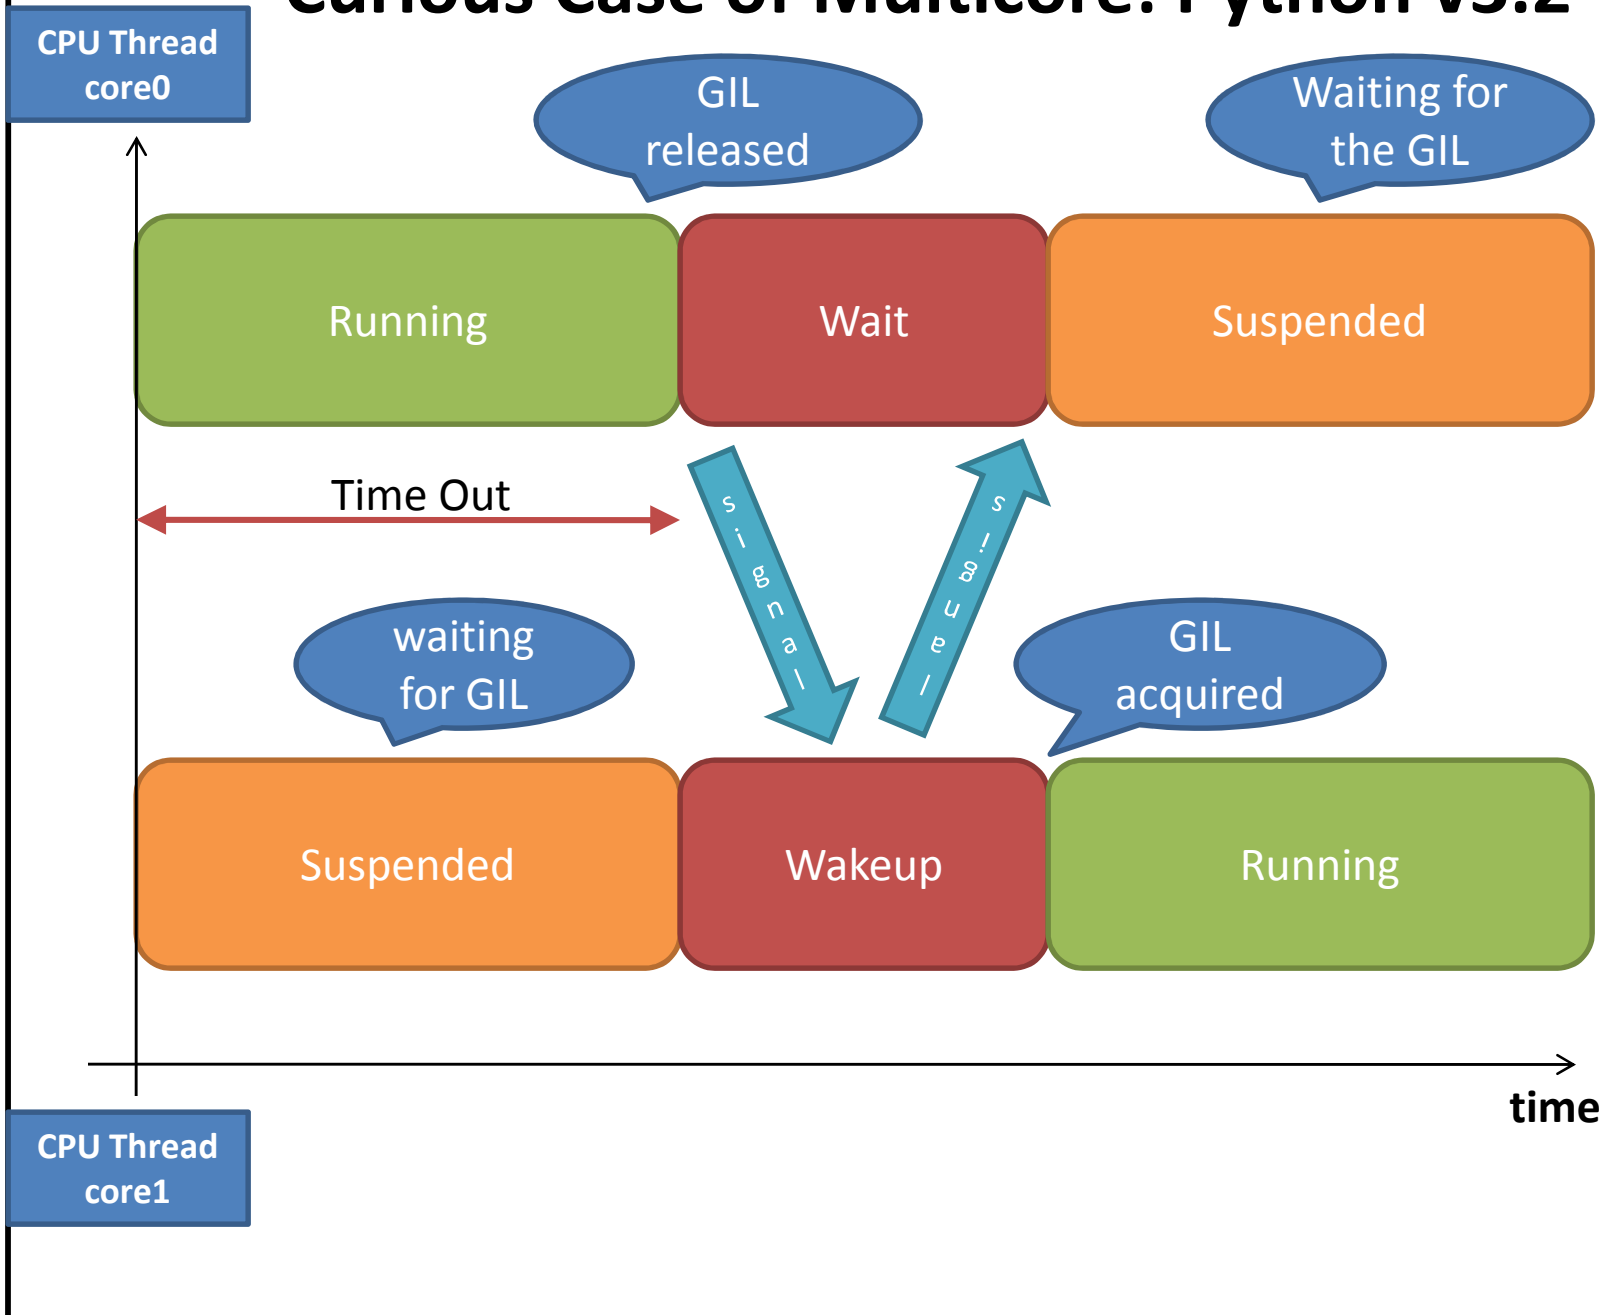
GIL



# New GIL: Python v3.2

- Regular “check” are discontinued
- We have new time-out mechanism.
  - Default time-out= 5ms
  - Configurable through `sys.setswitchinterval()`
- For every time-out, the current GIL holder is forced to release GIL
- It then signals the other waiting threads
- Waits for a signal from new GIL owner (acknowledgement).
- A sleeping thread wakes up, acquires the GIL, and signals the last owner.

# Curious Case of Multicore: Python v3.2





## Positive impact with new GIL

- Better GIL arbitration
  - Ensures that a thread runs only for 5ms
- Less context switching and fewer signals
- Multicore perspective: GIL battle eliminated!
- More responsive threads (fair scheduling)
- All iz well 😊

# I/O threads in Python

- An interesting optimization by interpreter
  - I/O calls are assumed blocking
- Python I/O extensively exercise this optimization with file, socket ops (e.g. read, write, send, recv calls)

*./Python3.2.1/Include/ceval.h*

*Py\_BEGIN\_ALLOW\_THREADS*

*Do some blocking I/O operation ...*

*Py\_END\_ALLOW\_THREADS*

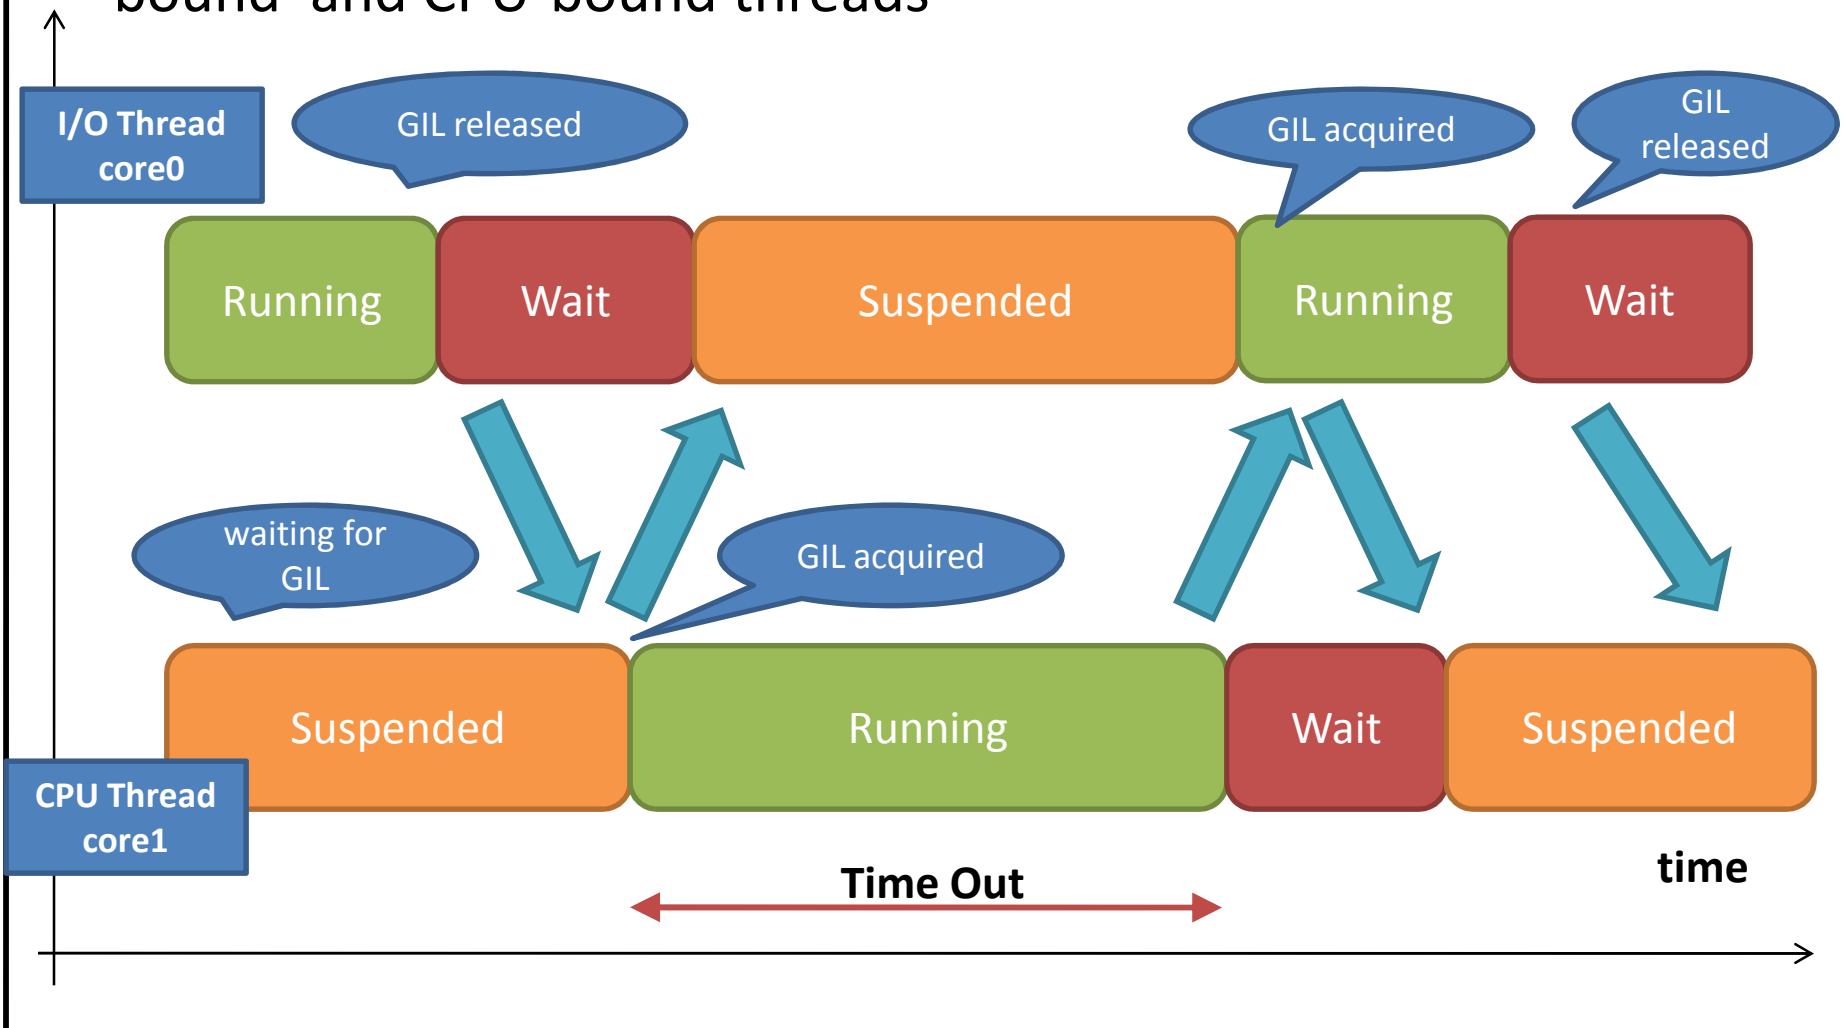
- I/O thread always releases the GIL

# Convoy effect: Fallout of I/O optimization

- When an I/O thread releases the GIL, another 'runnable' CPU bound thread can acquire it (remember we are on multiple cores).
- It leaves the I/O thread waiting for another time-out (default: 5ms)!
- Once CPU thread releases GIL, I/O thread acquires and releases it again
- This cycle goes on => performance suffers 😞

# Convoy “in” effect

**Convoy effect**- observed in an application comprising I/O-bound and CPU-bound threads



# Performance measurements!

- Curious to know how convoy effect translates into performance numbers
- We performed following tests with Python3.2:
  - An application with a CPU and a I/O thread
  - Executed on a dual core machine
  - CPU thread spends less than few seconds (<10s)!


I/O thread with CPU thread	I/O thread without CPU thread
97 seconds	23 seconds

# Comparing: Python 2.7 & Python 3.2


Python v2.7	Execution Time
Single Core	74 s
Dual Core	116 s

Python v3.2	Execution Time
Single Core	55 s
Dual Core	65 s

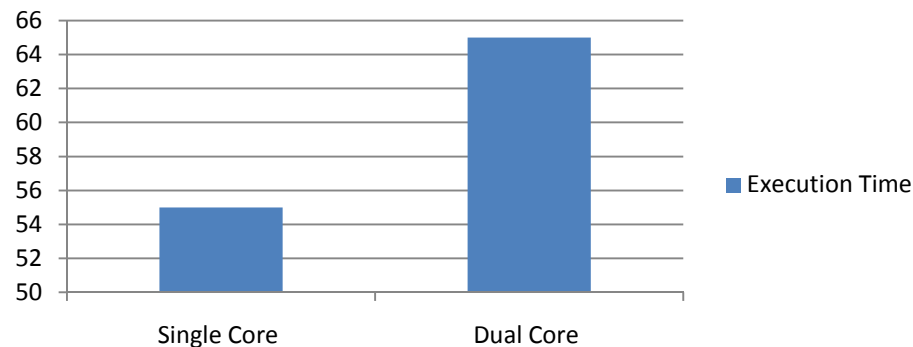
Execution Time – Single Core



Execution Time – Dual Core




Execution Time – Python v3.2



Performance  
dip still  
observed in  
dual cores ☹

# Getting into the problem space!!



Python  
v2.7 / v3.2

GIL

# Solution space!!

# GIL free world: Jython

- Jython is free of GIL 😊
- It can fully exploit multiple cores, as per our experiments
- Experiments with Jython2.5
  - Run with two CPU threads in tandem

Jython2.5	Execution time
Single core	44 s
Dual core	25 s

- Experiment shows performance improvement on a multi-core system

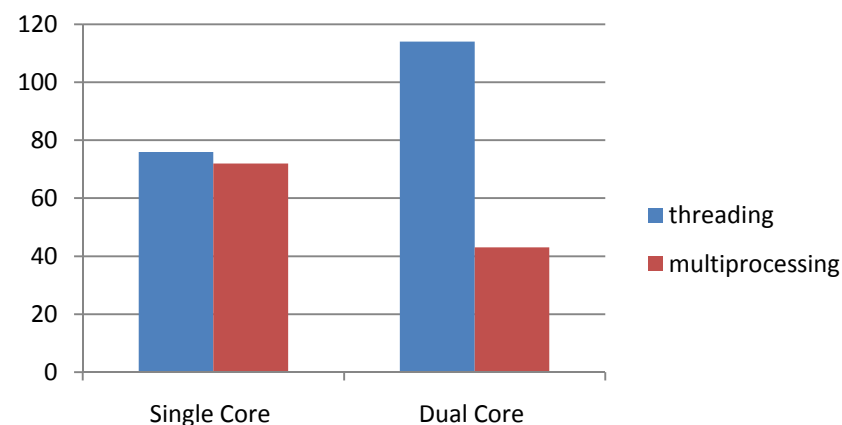


# Avoiding GIL impact with multiprocessing

- multiprocessing — Process-based “threading” interface
- “multiprocessing” module spawns a new Python interpreter instance for a process.
- Each process is independent, and GIL is irrelevant;
  - Utilizes multiple cores better than threads.
  - Shares API with “threading” module.

Python v2.7	Single Core	Dual Core
threading	76 s	114 s
multiprocessing	72 s	43 s

Cool! 40 % improvement in Execution Time on dual core!! 😊



# Conclusion

- Multi-core systems are becoming ubiquitous
- Python applications should exploit this abundant power
- CPython inherently suffers the GIL limitation
- An intelligent awareness of Python interpreter behavior is helpful in developing multi-threaded applications
- Understand and use 😊

# Questions

Thank you for your time and attention 😊

- Please share your feedback/ comments/ suggestions to us at:
- [cjgiridhar@gmail.com](mailto:cjgiridhar@gmail.com) , <http://technobeans.com>
- [vishalkanaujia@gmail.com](mailto:vishalkanaujia@gmail.com), <http://freethreads.wordpress.com>

# References

- Understanding the Python GIL, <http://dabeaz.com/talks.html>
- GlobalInterpreterLock, <http://wiki.python.org/moin/GlobalInterpreterLock>
- Thread State and the Global Interpreter Lock, <http://docs.python.org/c-api/init.html#threads>
- Python v3.2.2 and v2.7.2 documentation, <http://docs.python.org/>
- Concurrency and Python, <http://drdobbs.com/open-source/206103078?pgno=3>

Backup slides

# Python: GIL

- A thread needs GIL before updating Python objects, calling C/Python API functions
- Concurrency is emulated with regular 'checks' to switch threads
- Applicable to only CPU bound thread
- A blocking I/O operation implies relinquishing the GIL
  - `./Python2.7.5/Include/ceval.h`  
*Py\_BEGIN\_ALLOW\_THREADS*  
*Do some blocking I/O operation ...*  
*Py\_END\_ALLOW\_THREADS*
- Python file I/O extensively exercise this optimization

# GIL: Internals

- The function `Py_Initialize()` creates the GIL
- A thread create request in Python is just a `pthread_create()` call
- `../Python/ceval.c`
- `static PyThread_type_lock interpreter_lock = 0;`  
`/* This is the GIL */`
- o) `thread_PyThread_start_new_thread`: we call it for "each" user defined thread.
- `calls PyEval_InitThreads() ->`  
`PyThread_acquire_lock() {}`

# GIL: in action

- Each CPU bound thread requires GIL
- 'ticks count' determine duration of GIL hold
- `new_threadstate()` -> `tick_counter`
- We keep a list of Python threads and each thread-state has its `tick_counter` value
- As soon as tick decrements to zero, the thread release the GIL.



# GIL: Details

```
thread_PyThread_start_new_thread() ->  
void PyEval_InitThreads(void)  
{  
    if (interpreter_lock)  
        return;  
    interpreter_lock = PyThread_allocate_lock();  
    PyThread_acquire_lock(interpreter_lock, 1);  
    main_thread = PyThread_get_thread_ident();  
}
```

# Convoy effect: Python v2?

- Convoy effect holds true for Python v2 also
- The smaller interval of 'check' saves the day!
  - I/O threads don't have to wait for a longer time (5 m) for CPU threads to finish
  - Should choose the `setswitchinterval()` wisely
- The effect is not so visible in Python v2.0

# Stackless Python

- A different way of creating threads: Microthreads!
- No improvement from multi-core perspective
- Round-robin scheduling for “tasklets”
- Sequential execution ☹️